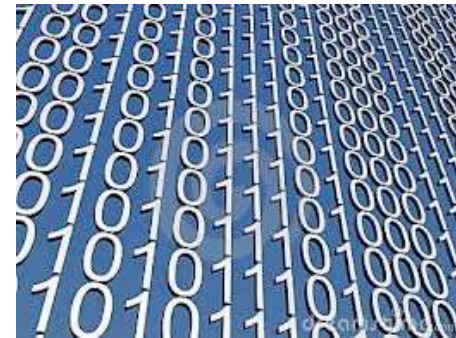


Number Systems and Number Representation



Goals of this Lecture

Help you learn (or refresh your memory) about:

- The binary, hexadecimal, and octal number systems
- Finite representation of unsigned integers
- Finite representation of signed integers
- Finite representation of rational numbers (if time)

Why?

- A power programmer must know number systems and data representation to fully understand C's **primitive data types**



Primitive values and
the operations on them

Agenda

Number Systems (Lecture 1)

Finite representation of unsigned integers (Lecture 2)

Finite representation of signed integers (Lecture 3)

Finite representation of rational numbers (Lecture 4)

The Decimal Number System

Name

- “decem” (Latin) => ten

Characteristics

- Ten symbols
 - 0 1 2 3 4 5 6 7 8 9
- Positional
 - $2945 \neq 2495$
 - $2945 = (2 \cdot 10^3) + (9 \cdot 10^2) + (4 \cdot 10^1) + (5 \cdot 10^0)$

(Most) people use the decimal number system

The Binary Number System

Name

- “binarius” (Latin) => two

Characteristics

- Two symbols
 - 0 1
- Positional
 - $1010_B \neq 1100_B$

Most (digital) computers use the binary number system

Terminology

- **Bit**: a binary digit
- **Byte**: (typically) 8 bits

Decimal-Binary Equivalence

<u>Decimal</u>	<u>Binary</u>
0	0
1	1
2	10
3	11
4	100
5	101
6	110
7	111
8	1000
9	1001
10	1010
11	1011
12	1100
13	1101
14	1110
15	1111

<u>Decimal</u>	<u>Binary</u>
16	10000
17	10001
18	10010
19	10011
20	10100
21	10101
22	10110
23	10111
24	11000
25	11001
26	11010
27	11011
28	11100
29	11101
30	11110
31	11111
...	...

Decimal-Binary Conversion

Binary to decimal: expand using positional notation

$$\begin{aligned} 100101_B &= (1*2^5) + (0*2^4) + (0*2^3) + (1*2^2) + (0*2^1) + (1*2^0) \\ &= 32 + 0 + 0 + 4 + 0 + 1 \\ &= 37 \end{aligned}$$

Decimal-Binary Conversion

Decimal to binary: do the reverse

- Determine largest power of $2 \leq \text{number}$; write template

$$37 = (? * 2^5) + (? * 2^4) + (? * 2^3) + (? * 2^2) + (? * 2^1) + (? * 2^0)$$

- Fill in template

$$37 = (1 * 2^5) + (0 * 2^4) + (0 * 2^3) + (1 * 2^2) + (0 * 2^1) + (1 * 2^0)$$

<u>-32</u>	
5	
<u>-4</u>	
1	
<u>-1</u>	
0	

100101_B

Decimal-Binary Conversion

Decimal to binary shortcut

- Repeatedly divide by 2, consider remainder

37	/	2	=	18	R	1
18	/	2	=	9	R	0
9	/	2	=	4	R	1
4	/	2	=	2	R	0
2	/	2	=	1	R	0
1	/	2	=	0	R	1



Read from bottom
to top: 100101_B

The Hexadecimal Number System

Name

- “hexa” (Greek) => six
- “decem” (Latin) => ten

Characteristics

- Sixteen symbols
 - 0 1 2 3 4 5 6 7 8 9 A B C D E F
- Positional
 - $A13D_H \neq 3DA1_H$

Computer programmers often use the hexadecimal number system

Decimal-Hexadecimal Equivalence

<u>Decimal</u>	<u>Hex</u>
0	0
1	1
2	2
3	3
4	4
5	5
6	6
7	7
8	8
9	9
10	A
11	B
12	C
13	D
14	E
15	F

<u>Decimal</u>	<u>Hex</u>
16	10
17	11
18	12
19	13
20	14
21	15
22	16
23	17
24	18
25	19
26	1A
27	1B
28	1C
29	1D
30	1E
31	1F

<u>Decimal</u>	<u>Hex</u>
32	20
33	21
34	22
35	23
36	24
37	25
38	26
39	27
40	28
41	29
42	2A
43	2B
44	2C
45	2D
46	2E
47	2F
...	...

Decimal-Hexadecimal Conversion

Hexadecimal to decimal: expand using positional notation

$$\begin{aligned} 25_{\text{H}} &= (2 \cdot 16^1) + (5 \cdot 16^0) \\ &= 32 + 5 \\ &= 37 \end{aligned}$$

Decimal to hexadecimal: use the shortcut

$$\begin{array}{l} 37 / 16 = 2 \text{ R } 5 \\ 2 / 16 = 0 \text{ R } 2 \end{array}$$



Read from bottom
to top: 25_{H}

Binary-Hexadecimal Conversion

Observation: $16^1 = 2^4$

- Every 1 hexadecimal digit corresponds to 4 binary digits

Binary to hexadecimal

1010	0001	0011	1101	_B
A	1	3	D	_H

Digit count in binary number
not a multiple of 4 =>
pad with zeros on left

Hexadecimal to binary

A	1	3	D	_H
1010	0001	0011	1101	_B

Discard leading zeros
from binary number if
appropriate

The Octal Number System

Name

- “octo” (Latin) => eight

Characteristics

- Eight symbols
 - 0 1 2 3 4 5 6 7
- Positional
 - $1743_8 \neq 7314_8$

Computer programmers often use the octal number system

Decimal-Octal Equivalence

<u>Decimal</u>	<u>Octal</u>
0	0
1	1
2	2
3	3
4	4
5	5
6	6
7	7
8	10
9	11
10	12
11	13
12	14
13	15
14	16
15	17

<u>Decimal</u>	<u>Octal</u>
16	20
17	21
18	22
19	23
20	24
21	25
22	26
23	27
24	30
25	31
26	32
27	33
28	34
29	35
30	36
31	37

<u>Decimal</u>	<u>Octal</u>
32	40
33	41
34	42
35	43
36	44
37	45
38	46
39	47
40	50
41	51
42	52
43	53
44	54
45	55
46	56
47	57
...	...

Decimal-Octal Conversion

Octal to decimal: expand using positional notation

$$\begin{aligned} 37_o &= (3 \cdot 8^1) + (7 \cdot 8^0) \\ &= 24 + 7 \\ &= 31 \end{aligned}$$

Decimal to octal: use the shortcut

$$\begin{array}{rcl} 31 & / & 8 = 3 \text{ R } 7 \\ 3 & / & 8 = 0 \text{ R } 3 \end{array}$$



Read from bottom
to top: 37_o

Binary-Octal Conversion

Observation: $8^1 = 2^3$

- Every 1 octal digit corresponds to 3 binary digits

Binary to octal

001	010	000	100	111	101	_B
1	2	0	4	7	5	_O

Digit count in binary number
not a multiple of 3 =>
pad with zeros on left

Octal to binary

1	2	0	4	7	5	_O
001	010	000	100	111	101	_B

Discard leading zeros
from binary number if
appropriate

Agenda

Number Systems (Lecture 1)

Finite representation of unsigned integers (Lecture 2)

Finite representation of signed integers (Lecture 3)

Finite representation of rational numbers (Lecture 4)

Bitwise Operations

Bitwise AND

- Similar to logical AND (& &), except it works on a bit-by-bit manner
- Denoted by a single ampersand: &

$$\begin{array}{r} (1001 \ \& \\ 0101) = \\ 0001 \end{array}$$

Bitwise OR

- Similar to logical OR (`||`), except it works on a bit-by-bit manner
- Denoted by a single pipe character: `|`

```
(1001 |  
 0101) =  
1101
```

Bitwise XOR

- Exclusive OR, denoted by a carat: ^
- Similar to bitwise OR, except that if both inputs are 1 then the result is 0

$$\begin{array}{r} (1001 \text{ } ^\wedge \\ 0101) = \\ 1100 \end{array}$$

Bitwise NOT

- Similar to logical NOT (!), except it works on a bit-by-bit manner
- Denoted by a tilde character: ~

$$\begin{array}{rcl} \sim 1001 & = & \\ & & 0110 \end{array}$$

Unsigned Data Types: Java vs. C

Java has type

- `int`
 - Can represent signed integers

C has type:

- `signed int`
 - Can represent signed integers
- `int`
 - Same as `signed int`
- `unsigned int`
 - Can represent only unsigned integers

To understand C, must consider representation of both unsigned and signed integers

Representing Unsigned Integers

Mathematics

- Range is 0 to ∞

Computer programming

- Range limited by computer's **word** size
- Word size is n bits \Rightarrow range is 0 to $2^n - 1$
- Exceed range \Rightarrow **overflow**

Nobel computers with gcc217

- $n = 32$, so range is 0 to $2^{32} - 1$ (4,294,967,295)

Pretend computer

- $n = 4$, so range is 0 to $2^4 - 1$ (15)

Hereafter, assume word size = 4

- All points generalize to word size = 32, word size = n

Representing Unsigned Integers

On pretend computer

<u>Unsigned Integer</u>	<u>Rep</u>
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111
8	1000
9	1001
10	1010
11	1011
12	1100
13	1101
14	1110
15	1111

Adding Unsigned Integers

Addition

		1
3		0011 _B
+ 10	+	1010 _B
--		----
13		1101 _B

Start at right column
Proceed leftward
Carry 1 when necessary

		11
7		0111 _B
+ 10	+	1010 _B
--		----
1		1 0001 _B

Beware of overflow

Results are mod 2^4

Subtracting Unsigned Integers

Subtraction

		12	
		0202	
10		1010 _B	
- 7	-	0111 _B	
--		----	
3		0011 _B	

Start at right column
Proceed leftward
Borrow 2 when necessary

		2	
3		0011 _B	
- 10	-	1010 _B	
--		----	
9		1001 _B	

Beware of overflow

Results are mod 2⁴

Shifting Unsigned Integers

Bitwise right shift (\gg): fill on left with zeros

10 \gg 1 \Rightarrow 5
 1010_B 0101_B

10 \gg 2 \Rightarrow 2
 1010_B 0010_B

What is the effect
arithmetically? (No
fair looking ahead)

Bitwise left shift (\ll): fill on right with zeros

5 \ll 1 \Rightarrow 10
 0101_B 1010_B

3 \ll 2 \Rightarrow 12
 0011_B 1100_B

What is the effect
arithmetically? (No
fair looking ahead)

Results are mod 2^4

Shift Left

- Move all the bits N positions to the left, subbing in N 0s on the right

Shift Left

- Move all the bits N positions to the left, subbing in N 0s on the right

1001

Shift Left

- Move all the bits N positions to the left, adding in N 0s on the right

$$\begin{array}{l} 1001 \ll 2 = \\ 100100 \end{array}$$

Shift Left

- Useful as a restricted form of multiplication
- Question: how?

$$\begin{array}{r} 1001 \ll 2 = \\ 100100 \end{array}$$

Shift Left as Multiplication

- Equivalent decimal operation:

234

Shift Left as Multiplication

- Equivalent decimal operation:

$$\begin{array}{r} 234 \ll 1 = \\ 2340 \end{array}$$

Shift Left as Multiplication

- Equivalent decimal operation:

$$\begin{array}{l} 234 \ll 1 = \\ 2340 \end{array}$$

$$\begin{array}{l} 234 \ll 2 = \\ 23400 \end{array}$$

Multiplication

- Shifting left N positions multiplies by $(base)^N$
- Multiplying by 2 or 4 is often necessary (shift left 1 or 2 positions, respectively)
- Often a whooole lot faster than telling the processor to multiply

$$234 \ll 2 = 23400$$

Shift Right

- Move all the bits N positions to the right, subbing in **either** N 0s or N 1s on the left
- Two different forms

Shift Right

- Move all the bits N positions to the right, subbing in **either** N 0s or N (whatever the leftmost bit is)s on the left
- Two different forms

$1001 \gg 2 =$
either 0010 **or** 1110

Shift Right as Division

- Question: If shifting left multiplies, what does shift right do?
 - Answer: divides in a similar way, but truncates result

Shift Right as Division

- Question: If shifting left multiplies, what does shift right do?
 - Answer: divides in a similar way, but truncates result

234

Shift Right as Division

- Question: If shifting left multiplies, what does shift right do?
 - Answer: divides in a similar way, but truncates result

$$\begin{array}{r} 234 \\ 23 \end{array} \gg 1 =$$

Other Operations on Unsigned Ints

Bitwise NOT (~)

- Flip each bit

$\sim 10 \Rightarrow 5$

$1010_B \quad 0101_B$

Bitwise AND (&)

- Logical AND corresponding bits

10		1010 _B
& 7		& 0111 _B
--		----
2		0010 _B

Useful for setting
selected bits to 0

Other Operations on Unsigned Ints

Bitwise OR: (|)

- Logical OR corresponding bits

10	1010 _B
1	0001 _B
--	----
11	1011 _B

Useful for setting
selected bits to 1

Bitwise exclusive OR (^)

- Logical exclusive OR corresponding bits

10	1010 _B
^ 10	^ 1010 _B
--	----
0	0000 _B

$x \wedge x$ sets
all bits to 0

The binary **XOR** operation will always produce a **1** output if either of its inputs is **1** and will produce a **0** output if both of its inputs are **0** or **1**.

Aside: Using Bitwise Ops for Arith

Can use \ll , \gg , and $\&$ to do some arithmetic efficiently

$$x * 2^y == x \ll y$$

$$\bullet 3 * 4 = 3 * 2^2 = 3 \ll 2 \Rightarrow 12$$

$0011_B \qquad 1100_B$

Fast way to **multiply**
by a power of 2

$$x / 2^y == x \gg y$$

$$\bullet 13 / 4 = 13 / 2^2 = 13 \gg 2 \Rightarrow 3$$

$1101_B \qquad 0011_B$

Fast way to **divide**
by a power of 2

$$x \% 2^y == x \& (2^y - 1)$$

$$\bullet 13 \% 4 = 13 \% 2^2 = 13 \& (2^2 - 1) \\ = 13 \& 3 \Rightarrow 1$$

Fast way to **mod**
by a power of 2

13	1101 _B
& 3	& 0011 _B
--	----
1	0001 _B

Two Forms of Shift Right

- Shifting in 0s makes sense
- What about shifting in the leftmost bit?
 - And why is this called “arithmetic” shift right?

1100 (arithmetic) >> 1 =
1110

Answer...Sort of

- Arithmetic form is intended for numbers in *two's complement*, whereas the non-arithmetic form is intended for *unsigned* numbers

Agenda

Number Systems (Lecture 1)

Finite representation of unsigned integers (Lecture 2)

Finite representation of signed integers (Lecture 3)

Finite representation of rational numbers (Lecture 4)

Signed Magnitude

<u>Integer</u>	<u>Rep</u>
-7	1111
-6	1110
-5	1101
-4	1100
-3	1011
-2	1010
-1	1001
-0	1000
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111

Definition

High-order bit indicates sign

0 => positive

1 => negative

Remaining bits indicate magnitude

$$1101_B = -101_B = -5$$

$$0101_B = 101_B = 5$$



Signed Magnitude (cont.)

<u>Integer</u>	<u>Rep</u>
-7	1111
-6	1110
-5	1101
-4	1100
-3	1011
-2	1010
-1	1001
-0	1000
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111

Computing negative

$\text{neg}(x) = \text{flip high order bit of } x$

$$\text{neg}(0101_B) = 1101_B$$

$$\text{neg}(1101_B) = 0101_B$$

Pros and cons

- + easy for people to understand
- + symmetric
- two reps of zero
- one of the bit patterns is wasted.
- addition doesn't work the way we want it to.

Signed Magnitude (cont.)

Problem #1: "The Case of the Missing Bit Pattern":

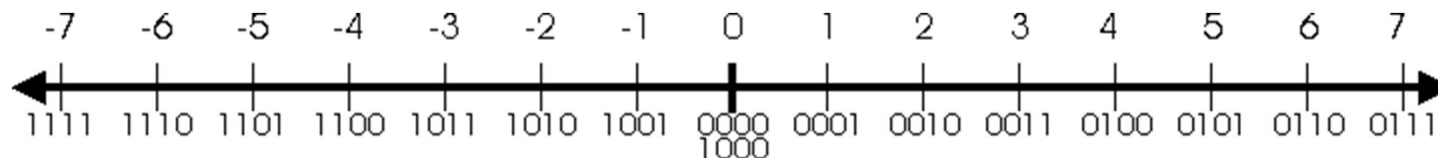
How many possible bit patterns can be created with 4 bits?

Easy, we know that's 16. In unsigned representation, we were able to represent 16 numbers: 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, and 15.

But with signed magnitude, we are only able to represent 15 numbers: -7, -6, -5, -4, -3, -2, -1, 0, 1, 2, 3, 4, 5, 6, and 7.

There's still 16 bit patterns, but one of them is either not being used or is duplicating a number. That bit pattern is '1000B'.

When we interpret this pattern, we get '-0' which is both nonsensical (negative zero? come on!) and redundant (we already have '0000B' to represent 0).



Signed Magnitude (cont.)

Problem #2: "Requires Special Care and Feeding": Remember we wanted to have negative binary numbers so we could use our binary addition algorithm to simulate binary subtraction. How does signed magnitude fare with addition? To test it, let's try subtracting 2 from 5 by adding 5 and -2. A positive 5 would be represented with the bit pattern '0101B' and -2 with '1010B'. Let's add these two numbers and see what the result is:

$$\begin{array}{r} 0101 \\ +1010 \\ \hline 1111 \end{array}$$

Now we interpret the result as a signed magnitude number. The sign is '1' (negative) and the magnitude is '7'. So the answer is a negative 7. But, wait a minute, $5-2=3$! This obviously didn't work.

Conclusion: signed magnitude doesn't work with regular binary addition algorithms.

Ones' Complement

<u>Integer</u>	<u>Rep</u>
-7	1000
-6	1001
-5	1010
-4	1011
-3	1100
-2	1101
-1	1110
-0	1111
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111

Definition

High-order bit has weight -7 ($-2^n + 1$)

$$\begin{aligned} 1010_B &= (1 * -7) + (0 * 4) + (1 * 2) + (0 * 1) \\ &= -5 \end{aligned}$$

$$\begin{aligned} 0010_B &= (0 * -7) + (0 * 4) + (1 * 2) + (0 * 1) \\ &= 2 \end{aligned}$$

Ones' Complement (cont.)

<u>Integer</u>	<u>Rep</u>
-7	1000
-6	1001
-5	1010
-4	1011
-3	1100
-2	1101
-1	1110
-0	1111
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110

Computing negative

$$\text{neg}(x) = \sim x$$

$$\text{neg}(0101_B) = 1010_B$$

$$\text{neg}(1010_B) = 0101_B$$

Computing negative (alternative)

$$\text{neg}(x) = 1111_B - x$$

$$\begin{aligned}\text{neg}(0101_B) &= 1111_B - 0101_B \\ &= 1010_B\end{aligned}$$

$$\begin{aligned}\text{neg}(1010_B) &= 1111_B - 1010_B \\ &= 0101_B\end{aligned}$$

Pros and cons

+ symmetric

- two reps of zero

Two's Complement

<u>Integer</u>	<u>Rep</u>
-8	1000
-7	1001
-6	1010
-5	1011
-4	1100
-3	1101
-2	1110
-1	1111
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111

Definition

High-order bit has weight -8 (-2^n)

$$\begin{aligned} 1010_B &= (1 * -8) + (0 * 4) + (1 * 2) + (0 * 1) \\ &= -6 \end{aligned}$$

$$\begin{aligned} 0010_B &= (0 * -8) + (0 * 4) + (1 * 2) + (0 * 1) \\ &= 2 \end{aligned}$$

Two's Complement (cont.)

<u>Integer</u>	<u>Rep</u>
-8	1000
-7	1001
-6	1010
-5	1011
-4	1100
-3	1101
-2	1110
-1	1111
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111

Computing negative

$$\text{neg}(x) = \sim x + 1$$

$$\text{neg}(x) = \text{onescomp}(x) + 1$$

$$\text{neg}(0101_B) = 1010_B + 1 = 1011_B$$

$$\text{neg}(1011_B) = 0100_B + 1 = 0101_B$$

Pros and cons

- not symmetric
- + one rep of zero

Two's Complement (cont.)

Almost all computers use two's complement to represent signed integers

Why?

- Arithmetic is easy
- Will become clear soon

Hereafter, assume two's complement representation of signed integers

Adding Signed Integers

pos + pos

		11
3	0011 _B	
+ 3	+ 0011 _B	
--	----	
6	0110 _B	

pos + pos (overflow)

		111
7	0111 _B	
+ 1	+ 0001 _B	
--	----	
-8	1000 _B	

pos + neg

		1111
3	0011 _B	
+ -1	+ 1111 _B	
--	----	
2	10010 _B	

neg + neg

		11
-3	1101 _B	
+ -2	+ 1110 _B	
--	----	
-5	11011 _B	

neg + neg (overflow)

		1 1
-6	1010 _B	
+ -5	+ 1011 _B	
--	----	
5	10101 _B	

Subtracting Signed Integers

Perform subtraction
with borrows

			1
			22
3		0011 _B	
- 4	-	0100 _B	
<hr/>			
-1		1111 _B	

-5		1011 _B	
- 2	-	0010 _B	
<hr/>			
-7		1001 _B	

or

Compute two's comp
and add

3		0011 _B	
+ -4	+	1100 _B	
<hr/>			
-1		1111 _B	

			111
-5		1011	
+ -2	+	1110	
<hr/>			
-7		11001	

Negating Signed Ints: Math

Question: Why does two's comp arithmetic work?

Answer: $[-b] \bmod 2^4 = [\text{twoscomp}(b)] \bmod 2^4$

$$\begin{aligned} & [-b] \bmod 2^4 \\ &= [2^4 - b] \bmod 2^4 \\ &= [2^4 - 1 - b + 1] \bmod 2^4 \\ &= [(2^4 - 1 - b) + 1] \bmod 2^4 \\ &= [\text{onescomp}(b) + 1] \bmod 2^4 \\ &= [\text{twoscomp}(b)] \bmod 2^4 \end{aligned}$$

See Bryant & O'Hallaron book for much more info

Subtracting Signed Ints: Math

And so:

$$[a - b] \bmod 2^4 = [a + \text{twoscomp}(b)] \bmod 2^4$$

$$\begin{aligned} & [a - b] \bmod 2^4 \\ &= [a + 2^4 - b] \bmod 2^4 \\ &= [a + 2^4 - 1 - b + 1] \bmod 2^4 \\ &= [a + (2^4 - 1 - b) + 1] \bmod 2^4 \\ &= [a + \text{onescomp}(b) + 1] \bmod 2^4 \\ &= [a + \text{twoscomp}(b)] \bmod 2^4 \end{aligned}$$

See Bryant & O'Hallaron book for much more info

Shifting Signed Integers

Bitwise (**logical/arithmetic**) left shift (<<): fill on right with zeros

$$\begin{array}{l} 3 \ll 1 \Rightarrow 6 \\ 0011_B \quad 0110_B \end{array}$$

$$\begin{array}{l} -3 \ll 1 \Rightarrow -6 \\ 1101_B \quad 1010_B \end{array}$$

Shift by n =
multiplying by 2^n

Bitwise **arithmetic** right shift: fill on left **with sign bit**

$$\begin{array}{l} 6 \gg 1 \Rightarrow 3 \\ 0110_B \quad 0011_B \end{array}$$

$$\begin{array}{l} -6 \gg 1 \Rightarrow -3 \\ 1010_B \quad 1101_B \end{array}$$

Shift by n = dividing by 2^n
and Round-floor

Results are mod 2^4

Shifting Signed Integers (cont.)

Bitwise **logical** right shift: fill on left **with zeros**

$$\boxed{6 \gg 1 \Rightarrow 3}$$

$0110_B \quad 0011_B$

$$\boxed{-6 \gg 1 \Rightarrow 5} \quad ?$$

$1010_B \quad 0101_B$

Right shift (\gg) could be logical or arithmetic

- Compiler designer decides
- **Logical** shift is ideal for unsigned binary numbers
- **Arithmetic** shift is ideal for signed two's complement binary numbers

Other Operations on Signed Ints

Bitwise NOT (~)

- Same as with unsigned ints

Bitwise AND (&)

- Same as with unsigned ints

Bitwise OR: (|)

- Same as with unsigned ints

Bitwise exclusive OR (^)

- Same as with unsigned ints

Agenda

Number Systems (Lecture 1)

Finite representation of unsigned integers (Lecture 2)

Finite representation of signed integers (Lecture 3)

Finite representation of rational numbers (Lecture 4)

Number Systems

- So far, we have studied the following **integer** number systems in computer
 - Unsigned numbers
 - Sign/magnitude numbers
 - Two's complement numbers
- What about **rational numbers**?
 - A **rational** number is one that can be expressed as the **ratio** of two integers
 - Infinite range and precision
 - For example, 2.5, -10.04, 0.75 etc

Rational Numbers

- Two common notations to represent rational numbers in computer
 - Fixed-point numbers
 - Floating-point numbers

Computer science

- Finite range and precision
- Approximate using **floating point** number
 - Binary point “floats” across bits

Fixed-Point Numbers

- Fixed point notation has an **implied binary point** between the integer and fraction bits
 - The binary point is not a part of the representation but is implied
 - Example:
 - Fixed-point representation of 6.75 using 4 integer bits and 4 fraction bits:

01101100

0110.1100

$$2^2 + 2^1 + 2^{-1} + 2^{-2} = 6.75$$

- The number of integer and fraction bits must be agreed upon by those generating and those reading the number
 - There is no way of knowing the existence of the binary point except through agreement of those people interpreting the number

Signed Fixed-Point Numbers

- As with whole numbers, negative fractional numbers can be represented in two ways
 - Sign/magnitude notation
 - Two's complement notation
- Example:
 - -2.375 using 8 bits (4 bits each to represent integer and fractional parts)
 - $2.375 = 0010 . 0110$
 - Sign/magnitude notation: 1010 0110
 - Two's complement notation:
 1. flip all the bits: 1101 1001
 2. add 1:
$$\begin{array}{r} 1101\ 1001 \\ + \qquad\quad 1 \\ \hline 1101\ 1010 \end{array}$$
- Addition and subtraction works easily in computer with 2's complement notation like integer addition and subtraction

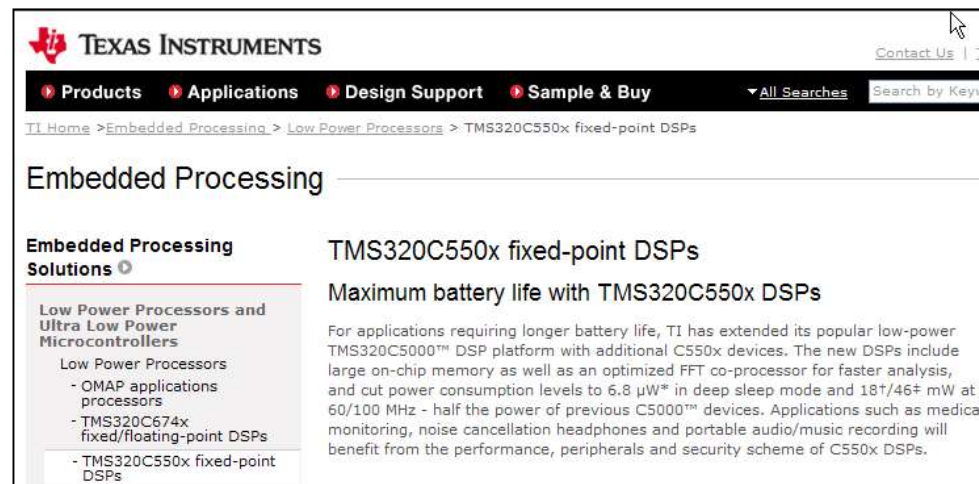
Example

- Suppose that we have 8 bits to represent a number
 - 4 bits for integer and 4 bits for fraction
- Compute $0.75 + (-0.625)$
 - $0.75 = 0000\ 1100$
 - $0.625 = 0000\ 1010$
 - -0.625 in 2's complement form: $1111\ 0110$

0.75	0000	1100
+ - 0.625	1111	0110
<hr/>		
0.125	0000	0010

Fixed-Point Number Systems

- Fixed-point number systems have a limitation of having a constant number of integer and fractional bits
- Some low-end digital signal processors support fixed-point numbers
 - Example: TMS320C550x TI (Texas Instruments) DSPs: www.ti.com



Floating-Point Numbers

- Floating-point number systems circumvent the limitation of having a constant number of integer and fractional bits
 - They allow the representation of very large and very small numbers
- The binary point **floats** to the right of the most significant 1
 - Similar to decimal scientific notation
 - For example, write 273_{10} in scientific notation:
 - Move the decimal point to the right of the most significant digit and increase the exponent:
$$273 = 2.73 \times 10^2$$

- In general, a number is written in scientific notation as:

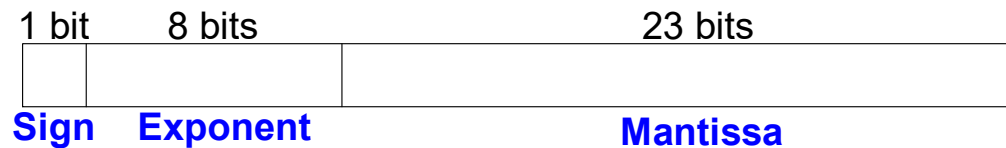
$$\pm M \times B^E$$

Where,

- M = mantissa
- B = base
- E = exponent
- In the example, M = 2.73, B = 10, and E = 2 (that is, $+2.73 \times 10^2$)

Floating-Point Numbers

- Floating-point number representation using 32 bits
 - 1 sign bit
 - 8 exponent bits
 - 23 bits for the mantissa.



- The following slides show **three** versions of floating-point representation with 228_{10} using a 32-bit
 - The final version is called the **IEEE 754 floating-point standard**

Floating-Point Representation #1

- First, convert the decimal number to binary
 - $228_{10} = 11100100_2 = 1.11001 \times 2^7$
- Next, fill in each field in the 32-bit:
 - The sign bit (1 bit) is positive, so 0
 - The exponent (8 bits) is 7 (111)
 - The mantissa (23 bits) is 1.11001

1 bit	8 bits	23 bits
0	00000111	11 1001 0000 0000 0000 0000
Sign	Exponent	Mantissa / Fraction

Floating-Point Representation #2

- You may have noticed that the first bit of the mantissa is always 1, since the binary point floats to the right of the most significant 1
 - Example: $228_{10} = 11100100_2 = \mathbf{1}.11001 \times 2^7$
- Thus, storing the most significant 1 (also called the implicit leading 1) is redundant information
- We can store just the **fraction parts** in the 23-bit field
 - Now, the leading 1 is implied

1 bit	8 bits	23 bits
0	0 0 0 0 0 1 1 1	110 0100 0000 0000 0000 0000
Sign	Exponent	Mantissa / Fraction

Floating-Point Representation #3

- The exponent needs to represent both positive and negative
- The final change is to use a **biased exponent**
 - The IEEE 754 standard uses a bias of 127
 - Biased exponent = bias + exponent
 - For example, an exponent of 7 is stored as $127 + 7 = 134 = 10000110_2$
- Thus , 228_{10} using the IEEE 754 32-bit floating-point standard is
 $228_{10} = 11100100_2 = 1.11001 \times 2^7$

1 bit	8 bits	23 bits
0	10000110	110 0100 0000 0000 0000 0000
Sign	Biased Exponent	Mantissa / Fraction

Most general purpose processors (including Intel and AMD processors) provide hardware support for double-precision floating-point numbers and operations

IEEE Floating Point Representation

Common finite representation: **IEEE floating point**

- More precisely: ISO/IEEE 754 standard

Using 32 bits (type float in C):

- 1 bit: sign (0=>positive, 1=>negative)
- 8 bits: exponent + 127
- 23 bits: binary fraction of the form 1.*dddddddddddddddddddddd*

Using 64 bits (type double in C):

- 1 bit: sign (0=>positive, 1=>negative)
- 11 bits: exponent + 1023
- 52 bits: binary fraction of the form
1.*dd*

1 bit	8 bits	23 bits
0	10000001	001 1000 0000 0000 0000 0000
Sign	Exponent	Mantissa / Fraction

Example

- Represent -58_{10} using the IEEE 754 floating-point standard
 - First, convert the decimal number to binary
 - $58_{10} = 111010_2 = 1.1101 \times 2^5$
 - Next, fill in each field in the 32-bit number
 - The sign bit is negative (1)
 - The 8 exponent bits are $(127 + 5) = 132 = 10000100_{(2)}$
 - The remaining 23 bits are the fraction bits: $11010000...000_{(2)}$

1 bit	8 bits	23 bits
1	10000100	110 1000 0000 0000 0000 0000
Sign	Exponent	Fraction

- It is 0xC2680000 in the hexadecimal form

Double Precision Example

- Represent -58_{10} using the IEEE 754 double precision
 - First, convert the decimal number to binary
 - $58_{10} = 111010_2 = 1.1101 \times 2^5$
 - Next, fill in each field in the 64-bit number
 - The sign bit is negative (1)
 - The 11 exponent bits are $(1023 + 5) = 1028 = 10000000100_{(2)}$
 - The remaining 52 bits are the fraction bits: $11010000...000_{(2)}$
 - It is 0xC04D000000000000 in the hexadecimal form

Floating-Point Numbers: Special Cases

- The IEEE 754 standard includes special cases for numbers that are difficult to represent, such as 0 because it lacks an implicit leading 1

Number	Sign	Exponent	Fraction
0	X	00000000	000000000000000000000000
∞	0	11111111	000000000000000000000000
$-\infty$	1	11111111	000000000000000000000000
NaN	X	11111111	non-zero

NaN is used for numbers that don't exist, such as $\sqrt{-1}$ or $\log(-5)$

Floating Point Example

Sign (1 bit):

- 1 => negative

11000001110110110000000000000000

32-bit representation

Exponent (8 bits):

- 10000011_B = 131
- 131 - 127 = 4

Fraction (23 bits):

- 1.10110110000000000000000_B
- 1 +
(1*2⁻¹) + (0*2⁻²) + (1*2⁻³) + (1*2⁻⁴) + (0*2⁻⁵) + (1*2⁻⁶) + (1*2⁻⁷)
= 1.7109375

Number:

- -1.7109375 * 2⁴ = -27.375

Floating Point Example 263.3

2	263	
2	131	1
2	65	1
2	32	1
2	16	0
2	8	0
2	4	0
2	2	0
2	1	0
	0	1

263: 100000111

IEEE754 floating-point standard can't represent some numbers exactly

0.3 * 2	0.6	0
0.6 * 2	1.2	1
0.2 * 2	0.4	0
0.4 * 2	0.8	0
0.8 * 2	1.6	1
0.6 * 2	1.2	1
		0
		0
		1
		1
		0

Stop when it gets 1.0

0.3 : 01001100110011....

Floating Point Example

1) 263.3

100000111.0100110011...

Sign (1 bit):

- positive \Rightarrow 0

Exponent (8 bits):

- $127 + 8 = 135$
- $135 = 10000111_B$

2) Scientific notation:

1.000001110100110011... * 2^8

└──────────────────┘
Mantissa

Fraction (23 bits):

- 00000111010011001100110

0100 0011 1000 0011 1010 0110 0110 0110

32-bit representation

Binary Coded Decimal (BCD)

- Since floating-point number systems can't represent some numbers exactly such as 0.3, some application (calculators) use BCD (Binary coded decimal)
 - BCD numbers encode each decimal digit using 4 bits with a range of 0 to 9

Decimal	BCD Digit
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111
8	1000
9	1001

BCD fixed-point notation examples

1.7 = 0001 . 0111

4.9 = 0100 . 1001

6.75 = 0110.01110101

- BCD is very common in electronic systems where a numeric value is to be displayed, especially, in systems consisting solely of digital logic (not containing a microprocessor) - Wiki

Converting Between Decimal and Binary Floating-Point Numbers

https://kyledewey.github.io/comp122-fall17/lecture/week_2/floating_point_interconversions.html

Summary

The binary, hexadecimal, and octal number systems

Finite representation of unsigned integers

Finite representation of signed integers

Finite representation of rational numbers

Essential for proper understanding of

- C primitive data types
- Assembly language
- Machine language

Backup Slides

Floating-Point Numbers: Rounding

- Arithmetic results that fall outside of the available precision must round to a neighboring number
- Rounding modes
 - Round down
 - Round up
 - Round toward zero
 - Round to nearest
- Example
 - Round 1.100101 (1.578125) so that it uses only 3 fraction bits
 - Round down: 1.100
 - Round up: 1.101
 - Round toward zero: 1.100
 - Round to nearest: 1.101
 - 1.625 is closer to 1.578125 than 1.5 is

Floating-Point Addition with the Same Sign

- Addition with floating-point numbers is not as simple as addition with 2's complement numbers
- The steps for adding floating-point numbers with the same sign are as follows
 1. Extract exponent and fraction bits
 2. Prepend leading 1 to form mantissa
 3. Compare exponents
 4. Shift smaller mantissa if necessary
 5. Add mantissas
 6. Normalize mantissa and adjust exponent if necessary
 7. Round result
 8. Assemble exponent and fraction back into floating-point format

Floating-Point Addition Example

Add the following floating-point numbers:

$$1.5 + 3.25$$

$$1.5_{(10)} = 1.1_{(2)} \times 2^0$$

$$3.25_{(10)} = 11.01_{(2)} = 1.101_{(2)} \times 2^1$$

$$1.1_{(10)} = 0x3FC00000 \text{ in IEEE 754 single precision}$$

$$3.25_{(10)} = 0x40500000 \text{ in IEEE 754 single precision}$$

Floating-Point Addition Example

1. Extract exponent and fraction bits

1 bit	8 bits	23 bits
0	01111111	100 0000 0000 0000 0000 0000
Sign	Exponent	Fraction

1 bit	8 bits	23 bits
0	10000000	101 0000 0000 0000 0000 0000
Sign	Exponent	Fraction

For first number (N1): $S = 0, E = 127, F = .1$

For second number (N2): $S = 0, E = 128, F = .101$

2. Prepend leading 1 to form mantissa

N1: 1.1

N2: 1.101

Floating-Point Addition Example

3. Compare exponents

$127 - 128 = -1$, so shift N1 right by 1 bit

4. Shift smaller mantissa if necessary

shift N1's mantissa: $1.1 \gg 1 = 0.11 \ (\times 2^1)$

5. Add mantissas

$$\begin{array}{r} 0.11 \times 2^1 \\ + 1.101 \times 2^1 \\ \hline 10.011 \times 2^1 \end{array}$$

Floating-Point Addition Example

6. Normalize mantissa and adjust exponent if necessary

$$10.011 \times 2^1 = 1.0011 \times 2^2$$

7. Round result

No need (fits in 23 bits)

8. Assemble exponent and fraction back into floating-point format

$$S = 0, E = 2 + 127 = 129 = 10000001_2, F = 001100..$$

1 bit	8 bits	23 bits
0	10000001	001 1000 0000 0000 0000 0000
Sign	Exponent	Fraction

$$4.75_{(10)} = \text{0x40980000} \text{ in the hexadecimal form}$$